

# HandlingOutliers

May 25, 2023

```
[238]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sbn
from sklearn import linear_model
from sklearn.model_selection import train_test_split # Sklearn package's
↳ randomized data splitting function
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import NearestNeighbors, KNeighborsClassifier,
↳ KNeighborsRegressor
```

```
[239]: import warnings
warnings.filterwarnings('ignore')
```

```
[240]: #import housing dataset and display first five rows
mhDataset = pd.read_csv("Melbourne_housing_FULL.csv")
mhDataset.head()
```

```
[240]:
```

	Suburb	Address	Rooms	Type	Price	Method	SellerG	\
0	Abbotsford	68 Studley St	2	h	NaN	SS	Jellis	
1	Abbotsford	85 Turner St	2	h	1480000.0	S	Biggin	
2	Abbotsford	25 Bloomburg St	2	h	1035000.0	S	Biggin	
3	Abbotsford	18/659 Victoria St	3	u	NaN	VB	Rounds	
4	Abbotsford	5 Charles St	3	h	1465000.0	SP	Biggin	

	Date	Distance	Postcode	...	Bathroom	Car	Landsize	BuildingArea	\
0	3/09/2016	2.5	3067.0	...	1.0	1.0	126.0	NaN	
1	3/12/2016	2.5	3067.0	...	1.0	1.0	202.0	NaN	
2	4/02/2016	2.5	3067.0	...	1.0	0.0	156.0	79.0	
3	4/02/2016	2.5	3067.0	...	2.0	1.0	0.0	NaN	
4	4/03/2017	2.5	3067.0	...	2.0	0.0	134.0	150.0	

	YearBuilt	CouncilArea	Lattitude	Longtitude	Regionname	\
0	NaN	Yarra City Council	-37.8014	144.9958	Northern Metropolitan	
1	NaN	Yarra City Council	-37.7996	144.9984	Northern Metropolitan	
2	1900.0	Yarra City Council	-37.8079	144.9934	Northern Metropolitan	
3	NaN	Yarra City Council	-37.8114	145.0116	Northern Metropolitan	
4	1900.0	Yarra City Council	-37.8093	144.9944	Northern Metropolitan	

```

Propertycount
0      4019.0
1      4019.0
2      4019.0
3      4019.0
4      4019.0

```

[5 rows x 21 columns]

```

[241]: cols_to_use = [
    ↳ ['Suburb', 'Rooms', 'Type', 'Price', 'Method', 'SellerG', 'Distance', 'Bedroom2', 'Bathroom', 'Car',
    ↳ 'CouncilArea', 'Regionname', 'Propertycount']
mhDataset = mhDataset[cols_to_use]
mhDataset.head()

```

```

[241]:      Suburb  Rooms Type      Price Method SellerG  Distance  Bedroom2  \
0  Abbotsford     2    h      NaN      SS   Jellis      2.5        2.0
1  Abbotsford     2    h  1480000.0      S   Biggin      2.5        2.0
2  Abbotsford     2    h  1035000.0      S   Biggin      2.5        2.0
3  Abbotsford     3    u      NaN      VB   Rounds      2.5        3.0
4  Abbotsford     3    h  1465000.0      SP   Biggin      2.5        3.0

```

```

      Bathroom  Car  Landsize  BuildingArea      CouncilArea  \
0          1.0  1.0    126.0          NaN  Yarra City Council
1          1.0  1.0    202.0          NaN  Yarra City Council
2          1.0  0.0    156.0      79.0  Yarra City Council
3          2.0  1.0      0.0          NaN  Yarra City Council
4          2.0  0.0    134.0    150.0  Yarra City Council

```

```

      Regionname  Propertycount
0  Northern Metropolitan    4019.0
1  Northern Metropolitan    4019.0
2  Northern Metropolitan    4019.0
3  Northern Metropolitan    4019.0
4  Northern Metropolitan    4019.0

```

```

[242]: mhDataset.isna().sum()/len(mhDataset)*100

```

```

[242]: Suburb      0.000000
Rooms      0.000000
Type       0.000000
Price     21.832057
Method     0.000000
SellerG    0.000000
Distance   0.002869
Bedroom2   23.573457

```

```

Bathroom      23.599277
Car            25.039447
Landsize       33.881286
BuildingArea   60.576068
CouncilArea    0.008607
Regionname     0.008607
Propertycount  0.008607
dtype: float64

```

```
[243]: missing_values_count = mhDataset.isnull().sum()
missing_values_count
```

```

[243]: Suburb      0
Rooms      0
Type       0
Price      7610
Method     0
SellerG    0
Distance   1
Bedroom2   8217
Bathroom   8226
Car        8728
Landsize   11810
BuildingArea 21115
CouncilArea 3
Regionname 3
Propertycount 3
dtype: int64

```

```
[244]: mhDataset = mhDataset.dropna()
mhDataset.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 9244 entries, 2 to 34856
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype
---  -
0   Suburb      9244 non-null   object
1   Rooms       9244 non-null   int64
2   Type        9244 non-null   object
3   Price       9244 non-null   float64
4   Method      9244 non-null   object
5   SellerG     9244 non-null   object
6   Distance    9244 non-null   float64
7   Bedroom2    9244 non-null   float64
8   Bathroom    9244 non-null   float64
9   Car         9244 non-null   float64
10  Landsize    9244 non-null   float64

```

```

11 BuildingArea  9244 non-null  float64
12 CouncilArea  9244 non-null  object
13 Regionname    9244 non-null  object
14 Propertycount 9244 non-null  float64
dtypes: float64(8), int64(1), object(6)
memory usage: 1.1+ MB

```

```
[245]: mhDataset.isna().sum()/len(mhDataset)*100
```

```

[245]: Suburb          0.0
Rooms              0.0
Type              0.0
Price             0.0
Method            0.0
SellerG           0.0
Distance          0.0
Bedroom2          0.0
Bathroom          0.0
Car               0.0
Landsize          0.0
BuildingArea      0.0
CouncilArea       0.0
Regionname        0.0
Propertycount     0.0
dtype: float64

```

```
[246]: mhDataset.isna().sum()
```

```

[246]: Suburb          0
Rooms              0
Type              0
Price             0
Method            0
SellerG           0
Distance          0
Bedroom2          0
Bathroom          0
Car               0
Landsize          0
BuildingArea      0
CouncilArea       0
Regionname        0
Propertycount     0
dtype: int64

```

```

[247]: mhDataset['Type'] = mhDataset['Type'].replace({'h': 'House/Villa', 'u': 'Unit/
↳Duplex', 't': 'TownHouse'})
mhDataset.head(10)

```

```

mhDataset['Method'] = mhDataset['Method'].replace({'SS':'Sold after auction_
↳price not disclosed',

                                                    'S':'Property Sold',
                                                    'VB':'Vendor Bid',
                                                    'SP':'Property Sold Prior',
                                                    'PI':'Property passed in',
                                                    'SN':'Sold not disclosed',
                                                    'W':'Withdrawn Prior',
                                                    'PN':'Sold prior not_
↳disclosed',

                                                    'SA':'Sold after auction'})

```

```
[248]: mhDataset.head()
```

```

[248]:
      Suburb  Rooms      Type      Price      Method SellerG \
2  Abbotsford      2  House/Villa  1035000.0      Property Sold  Biggin
4  Abbotsford      3  House/Villa  1465000.0  Property Sold Prior  Biggin
6  Abbotsford      4  House/Villa  1600000.0      Vendor Bid  Nelson
11 Abbotsford      3  House/Villa  1876000.0      Property Sold  Nelson
14 Abbotsford      2  House/Villa  1636000.0      Property Sold  Nelson

      Distance  Bedroom2  Bathroom  Car  Landsize  BuildingArea \
2           2.5         2.0         1.0  0.0     156.0          79.0
4           2.5         3.0         2.0  0.0     134.0         150.0
6           2.5         3.0         1.0  2.0     120.0         142.0
11          2.5         4.0         2.0  0.0     245.0         210.0
14          2.5         2.0         1.0  2.0     256.0         107.0

      CouncilArea      Regionname  Propertycount
2  Yarra City Council  Northern Metropolitan      4019.0
4  Yarra City Council  Northern Metropolitan      4019.0
6  Yarra City Council  Northern Metropolitan      4019.0
11 Yarra City Council  Northern Metropolitan      4019.0
14 Yarra City Council  Northern Metropolitan      4019.0

```

```
[249]: mhDataset.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 9244 entries, 2 to 34856
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Suburb      9244 non-null   object
1   Rooms       9244 non-null   int64
2   Type        9244 non-null   object
3   Price       9244 non-null   float64
4   Method      9244 non-null   object

```


```

5  SellerG      9244 non-null  object
6  Distance     9244 non-null  float64
7  Bedroom2     9244 non-null  float64
8  Bathroom     9244 non-null  float64
9  Car          9244 non-null  float64
10 Landsize     9244 non-null  float64
11 BuildingArea 9244 non-null  float64
12 CouncilArea  9244 non-null  object
13 Regionname   9244 non-null  object
14 Propertycount 9244 non-null  float64
dtypes: float64(8), int64(1), object(6)
memory usage: 1.1+ MB

```

```
[250]: mhDataset['Price'].plot(kind='kde')
```

```
[250]: <AxesSubplot:ylabel='Density'>
```



```
[251]: total_cells = np.product(mhDataset.shape)
total_missing = missing_values_count.sum()
```

```
# percent of data that is missing
(total_missing/total_cells) * 100
```

[251]: 47.39362469349488

[252]: mhDataset.describe().T

[252]:

	count	mean	std	min	25%	\
Rooms	9244.0	3.098118e+00	0.964029	1.0	2.0	
Price	9244.0	1.092329e+06	679621.207086	131000.0	641000.0	
Distance	9244.0	1.124115e+01	6.882570	0.0	6.4	
Bedroom2	9244.0	3.077347e+00	0.966366	0.0	2.0	
Bathroom	9244.0	1.652423e+00	0.724991	1.0	1.0	
Car	9244.0	1.695370e+00	0.975529	0.0	1.0	
Landsize	9244.0	5.288338e+02	1212.965090	0.0	210.0	
BuildingArea	9244.0	1.569946e+02	480.976260	0.0	100.0	
Propertycount	9244.0	7.463867e+03	4369.422310	249.0	4380.0	

	50%	75%	max
Rooms	3.0	4.0	12.0
Price	900000.0	1341250.0	9000000.0
Distance	10.3	13.9	48.1
Bedroom2	3.0	4.0	12.0
Bathroom	2.0	2.0	9.0
Car	2.0	2.0	10.0
Landsize	474.0	651.0	44500.0
BuildingArea	132.0	181.0	44515.0
Propertycount	6543.0	10331.0	21650.0

Looking at the table description, the values for Rooms and Bedroom2 looks almost equal. And also looking at the column names, we can either have Rooms column or Bedroom column. In this case, let us have “Rooms” column and delete the “Bedroom2” column

[253]: mhDataset.dropna(axis=1,inplace=True)

[254]: mhDataset = mhDataset.drop(['Bedroom2'],axis =1)

[255]: mhDataset.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9244 entries, 2 to 34856
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Suburb      9244 non-null   object
1   Rooms       9244 non-null   int64
2   Type        9244 non-null   object
3   Price       9244 non-null   float64
4   Method      9244 non-null   object
```

```

5  SellerG          9244 non-null  object
6  Distance         9244 non-null  float64
7  Bathroom         9244 non-null  float64
8  Car              9244 non-null  float64
9  Landsize         9244 non-null  float64
10 BuildingArea     9244 non-null  float64
11 CouncilArea      9244 non-null  object
12 Regionname       9244 non-null  object
13 Propertycount    9244 non-null  float64

```

dtypes: float64(7), int64(1), object(6)

memory usage: 1.1+ MB

Handling outliers using IQR method

The formula to exclude outliers using InterQuartile method - Dependent variable column  $> (\text{Quartile1} - (1.5 * \text{IQR}))$  and Dependent variable column  $\leq (\text{Quartile3} + (1.5 * \text{IQR}))$  Lower Limit =  $\text{Quartile1} - (1.5 * \text{IQR})$  Upper Limit =  $\text{Quartile3} + (1.5 * \text{IQR})$

```

[256]: #Now lets handle the outliers
mhsDataset = mhDataset.copy()

IQR = mhDataset['Price'].quantile(0.75) - mhDataset['Price'].quantile(0.25)

lower = mhDataset['Price'].quantile(0.25) - 1.5* IQR
upper = mhDataset['Price'].quantile(0.75) + 1.5* IQR

outliers = np.where(mhDataset['Price']>upper,True, np.
    ↳where(mhDataset['Price']<lower,True,False))

mhsDataset = mhDataset.loc[~(outliers)]

print("Lower Limit :",lower)
print("Uppwe Limit :",upper)
#mhsDataset = mhsDataset[~((mhsDataset['Price']<lower) &
    ↳(mhsDataset['Price']>upper))]

```

Lower Limit : -409375.0

Uppwe Limit : 2391625.0

```

[257]: mhDataset['Price'].plot(kind='kde')


```

```

[257]: <AxesSubplot:ylabel='Density'>

```






```
[258]: mhsDataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8788 entries, 2 to 34856
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Suburb          8788 non-null   object
1   Rooms           8788 non-null   int64
2   Type            8788 non-null   object
3   Price           8788 non-null   float64
4   Method          8788 non-null   object
5   SellerG         8788 non-null   object
6   Distance        8788 non-null   float64
7   Bathroom        8788 non-null   float64
8   Car             8788 non-null   float64
9   Landsize        8788 non-null   float64
10  BuildingArea    8788 non-null   float64
11  CouncilArea     8788 non-null   object
12  Regionname      8788 non-null   object
13  Propertycount   8788 non-null   float64
```

```
dtypes: float64(7), int64(1), object(6)
memory usage: 1.0+ MB
```

```
[259]: mhsDataset['Price'].plot(kind='kde')
```

```
[259]: <AxesSubplot:ylabel='Density'>
```



```
[260]: mhsDataset.describe().T
```

```
[260]:
```

	count	mean	std	min	25%	\
Rooms	8788.0	3.039713	0.935446	1.0	2.00	
Price	8788.0	985448.177856	464712.324595	131000.0	630000.00	
Distance	8788.0	11.412062	6.977235	0.0	6.50	
Bathroom	8788.0	1.599681	0.671289	1.0	1.00	
Car	8788.0	1.665794	0.959368	0.0	1.00	
Landsize	8788.0	517.113792	1237.520726	0.0	199.75	
BuildingArea	8788.0	150.325378	490.774992	0.0	98.00	
Propertycount	8788.0	7466.414998	4422.746339	249.0	4294.00	
		50%	75%	max		
Rooms		3.0	4.0	12.0		

Price	870000.0	1266000.0	2385000.0
Distance	10.5	14.0	48.1
Bathroom	2.0	2.0	9.0
Car	2.0	2.0	10.0
Landsize	454.5	643.0	44500.0
BuildingArea	130.0	174.0	44515.0
Propertycount	6543.0	10331.0	21650.0

```
[261]: independentCols = ['Rooms', 'Distance', 'Bathroom', 'Car', 'Landsize',
    ↪ 'BuildingArea', 'Propertycount']
xs =mhsDataset[independentCols]
ys=mhsDataset['Price']
```

```
[262]: Xs_train, Xs_test, Ys_train, Ys_test = train_test_split(xs, ys, test_size=0.30,
    ↪ random_state=100)
```

```
[263]: linearReg = LinearRegression()
linearReg.fit(Xs_train,Ys_train)
```

```
[263]: LinearRegression()
```

```
[264]: linearReg.score(Xs_train,Ys_train)
```

```
[264]: 0.3834886371405757
```

```
[265]: linearReg.score(Xs_test,Ys_test)
```

```
[265]: 0.37114890123584865
```

```
[266]: lassoReg = linear_model.Lasso(alpha=50,max_iter=100,tol=1)
lassoReg.fit(Xs_train,Ys_train)
```

```
[266]: Lasso(alpha=50, max_iter=100, tol=1)
```

```
[267]: lassoReg.score(Xs_train,Ys_train)
```

```
[267]: 0.38307425605552137
```

```
[268]: lassoReg.score(Xs_test,Ys_test)
```

```
[268]: 0.3691530493706163
```

```
[269]: ridgeReg = linear_model.Ridge(alpha=100,max_iter=999,tol=1)
ridgeReg.fit(Xs_train,Ys_train)
```

```
[269]: Ridge(alpha=100, max_iter=999, tol=1)
```

```
[270]: ridgeReg.score(Xs_test,Ys_test)
```

```
[270]: 0.37073499653704556
```

```
[271]: from sklearn.metrics import mean_squared_error, r2_score
import numpy as np
```

```
knn_model = KNeighborsRegressor().fit(Xs_train, Ys_train)
predicted_values = knn_model.predict(Xs_test)
```

```
[272]: predict_df = pd.DataFrame({"Dependent_Test" : Ys_test, "Dependent_Predicted" :
    ↪ predicted_values})
predict_df.head()
```

```
[272]:
```

	Dependent_Test	Dependent_Predicted
26064	1100000.0	1309000.0
10977	1287000.0	1558700.0
14051	490000.0	497200.0
8529	767500.0	768400.0
26212	1230000.0	1832000.0

```
[273]: predict_df = (predict_df*(np.max(mhsDataset.Price) - np.min(mhsDataset.Price)))
    ↪ + np.min(mhsDataset.Price)
```

```
[274]: from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

print("Mean Squared Error = ", mean_squared_error(predict_df.
    ↪ Dependent_Predicted, predict_df.Dependent_Test))

print("R2 Score", r2_score(predict_df.Dependent_Predicted, predict_df.
    ↪ Dependent_Test))
```

Mean Squared Error = 5.070665377882707e+23

R2 Score 0.2901374317807399

Handling Outliers using Normal Distribution

The formula to handle outliers using normal distribution is Dependent variable column  $>$  mean - 3 \* Standard Deviation and Dependent variable column  $\leq$  mean - 3 \* Standard Deviation Lower Limit = mean - 3 \* Standard Deviation Upper Limit = mean + 3 \* Standard Deviation

```
[275]: #Now lets handle the outliers
mhDatasetND = mhDataset.copy()

#IQR = mhDataset['Price'].quantile(0.75) - mhDataset['Price'].quantile(0.25)

lower = mhDataset['Price'].mean() - 3 * mhDataset['Price'].std()
upper = mhDataset['Price'].mean() + 3 * mhDataset['Price'].std()
```

```

outliers = np.where(mhDataset['Price']>upper,True, np.
    ↪where(mhDataset['Price']<lower,True,False))

mhDatasetND = mhDataset.loc[~(outliers)]

print("Lower Limit :",lower)
print("Upper Limit :",upper)
#mhDataset = mhDataset[~((mhDataset['Price']<lower) &
    ↪(mhDataset['Price']>upper))]


```

Lower Limit : -946535.0324432228

Upper Limit : 3131192.210071955


```
[276]: mhDataset['Price'].plot(kind='kde')
```

[276]: <AxesSubplot:ylabel='Density'>



```
[277]: mhDatasetND['Price'].plot(kind='kde')
```

[277]: <AxesSubplot:ylabel='Density'>



Looking at the above two graphs, the exclusion of outliers is evident.

```
[278]: mhDatasetND.describe().T
```

```
[278]:
```

	count	mean	std	min	25%	\
Rooms	9078.0	3.075677e+00	0.952607	1.0	2.0	
Price	9078.0	1.039889e+06	547912.433123	131000.0	636125.0	
Distance	9078.0	1.130724e+01	6.914404	0.0	6.4	
Bathroom	9078.0	1.628663e+00	0.695856	1.0	1.0	
Car	9078.0	1.679445e+00	0.963541	0.0	1.0	
Landsize	9078.0	5.235416e+02	1221.791081	0.0	207.0	
BuildingArea	9078.0	1.541873e+02	484.691677	0.0	99.0	
Propertycount	9078.0	7.465040e+03	4389.801336	249.0	4380.0	

	50%	75%	max
Rooms	3.0	4.0	12.0
Price	886000.0	1310000.0	3120000.0
Distance	10.4	14.0	48.1
Bathroom	2.0	2.0	9.0
Car	2.0	2.0	10.0
Landsize	465.0	650.0	44500.0

BuildingArea	131.0	179.0	44515.0
Propertycount	6543.0	10331.0	21650.0

```
[279]: independentCols = ['Rooms', 'Distance', 'Bathroom', 'Car', 'Landsize',
    ↪ 'BuildingArea', 'Propertycount']
xsn =mhDatasetND[independentCols]
ysn =mhDatasetND['Price']
```

```
[280]: Xsn_train, Xsn_test, Ysn_train, Ysn_test = train_test_split(xsn, ysn,
    ↪ test_size=0.30, random_state=100)
```

```
[281]: linearReg = LinearRegression()
linearReg.fit(Xsn_train,Ysn_train)
```

```
[281]: LinearRegression()
```

```
[282]: linearReg.score(Xsn_train,Ysn_train)
```

```
[282]: 0.4243021032287464
```

```
[283]: linearReg.score(Xs_test,Ys_test)
```

```
[283]: 0.3577993841828496
```

```
[284]: lassoReg = linear_model.Lasso(alpha=50,max_iter=100,tol=1)
lassoReg.fit(Xsn_train,Ysn_train)
```

```
[284]: Lasso(alpha=50, max_iter=100, tol=1)
```

```
[285]: lassoReg.score(Xsn_train,Ysn_train)
```

```
[285]: 0.42420586222745427
```

```
[286]: lassoReg.score(Xsn_test,Ysn_test)
```

```
[286]: 0.37889064944333106
```

```
[287]: ridgeReg = linear_model.Ridge(alpha=100,max_iter=999,tol=1)
ridgeReg.fit(Xsn_train,Ysn_train)
```

```
[287]: Ridge(alpha=100, max_iter=999, tol=1)
```

```
[288]: knn_model = KNeighborsRegressor().fit(Xsn_train, Ysn_train)
predicted_values = knn_model.predict(Xsn_test)
```

```
[289]: predict_df = pd.DataFrame({"Dependent_Test" : Ysn_test, "Dependent_Predicted" :
    ↪ predicted_values})
predict_df.head()
```

```
[289]:
```

	Dependent_Test	Dependent_Predicted
14949	975000.0	1053000.0
20120	1210000.0	1378300.0
34177	760000.0	866500.0
29263	910000.0	959200.0
6425	1120000.0	903800.0

```
[290]: predict_df = (predict_df*(np.max(mhDatasetND.Price) - np.min(mhDatasetND.
↪Price))) + np.min(mhDatasetND.Price)
```

```
[291]: print("Mean Squared Error = ", mean_squared_error(predict_df.
↪Dependent_Predicted, predict_df.Dependent_Test))

print("R2 Score",r2_score(predict_df.Dependent_Predicted,predict_df.
↪Dependent_Test))
```

Mean Squared Error = 1.1996829183761369e+24

R2 Score 0.2790573905413596

Interquartile Range Method Mean Squared Error = 5.070665377882707e+23 R2 Score 0.2901374317807399

Z-Score Method Mean Squared Error = 1.1996829183761369e+24 R2 Score 0.2790573905413596

Looking at the MSE and R2 score of both the methods, Z-Score method suits this dataset to handle outliers since there seems to be less error.

```
[ ]:
```